

Lab Tutorial 1

Course Project Part 1

Jean-Christophe Petkovich
j2petkov@uwaterloo.ca

University of Waterloo

January 15, 2014

Concepts

Memory Management

Process States

Scheduling

Initialization

Concepts in Lab Part 1

- ▶ Memory management systems
- ▶ Process states/transitions
- ▶ Context Switching
- ▶ Priority based scheduling
- ▶ Operating System (OS) initialization

Memory Management

You will be implementing a simple memory management system.


- ▶ Memory allocation will primarily be used for allocating blocks of memory to hold messages.
- ▶ It will be much easier in the long run to implement a generic memory allocation system.

Your allocation Application Programming Interface (API) will be analogous to malloc/free, it will have two methods:

- ▶ One for asking the OS for free memory blocks
- ▶ One for telling the OS you're finished with a block, so it can be returned to the heap

Memory Model

The simplest implementation of a heap is a linked list of memory blocks¹. Here is what your datastructure might look like if you use linked lists.



How do I turn a memory block into a useful thing? Casting! If the current block's pointer is NULL, you're out of blocks!

¹But there are many possible implementations.

Some notes...

A memory block like the previous one might look like this:

```
typedef struct mem_blk {  
    uint32_t *next_blk;  
}
```

- ▶ Note the lack of a mention of the memory region.
- ▶ Size and number of memory blocks should be configured with global constants, not by the struct.
- ▶ The list of free blocks is created during the initialization phase.

Request Memory Block Pseudo Code

```
int k_request_memory_block() {  
    atomic(on);  
    while (no memory block is available) {  
        put PCB on blocked_resource_q;  
        set process state to BLOCKED_ON_RESOURCE;  
        release_processor();  
    }  
    int mem_blk = next free block;  
    update the heap;  
    atomic(off);  
    return mem_blk;  
}
```

Hint: Create your own generic utilities for handling linked lists and queues, you will need them for other things.

Release Memory Block Pseudo Code

```
int k_release_memory_block(void* memory_block) {  
    atomic(on);  
    if (memory block pointer is not valid)  
        return ERROR_CODE;  
    put memory_block into heap;  
    if (blocked on resource q not empty) {  
        handle_process_ready(pop(blocked resource q));  
    }  
    atomic(off);  
    return SUCCESS_CODE;  
}
```

Hint: Create your own generic utilities for handling linked lists and queues, you will need them for other things.

Process Control Blocks

Each process will have a Process Control Block (PCB), this stores information about the process's status including things such as the:

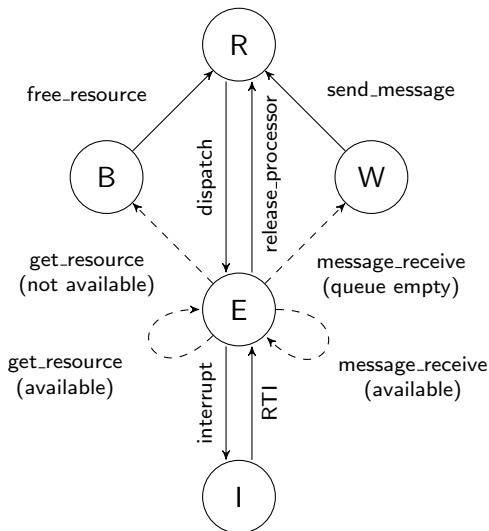
- ▶ Process's state
- ▶ ID
- ▶ Priority
- ▶ Program Counter (PC)
- ▶ CPU Registers
- ▶ Stack Pointer (SP)

State is what we're going to deal with next.

Process State Transition Diagram

R = ready, B = blocked on resource, W = waiting for message, E = executing, I = interrupted.

→ = forced, - -> = voluntary



Process Switching

Process switching:

1. Select next process to execute using scheduler
2. Invoke context switch to new process

Context switching²:

1. Save context of currently executing process
2. Change the process's state back to ready
3. Update `current_process` to new process
4. Set state of new process to executing
5. Restore context of `current_process`
6. Execute `current_process`

²You will get (have gotten?) *intentionally insufficient* sample code for context switching.

Scheduling

Your OS must have a scheduling algorithm that meets the following requirements:

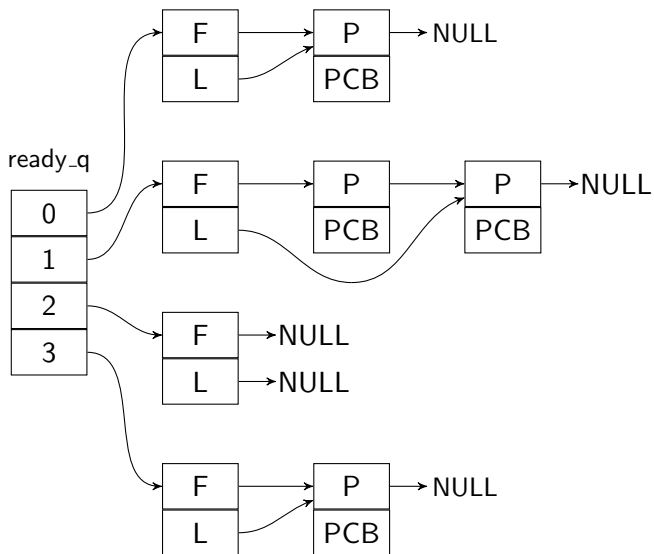
- ▶ Fixed, priority-based scheduling
- ▶ Each process has assigned priority
 - ▶ Highest priority process executes first
 - ▶ First-come-serve for processes of same priority

Scheduling Procedure

Procedure:

1. `process_switch` invokes scheduler
2. Scheduler selects highest priority ready process
3. `context_switch(new_proc)` lets the selected process execute

Scheduling Ready Queues



Note: 0 is traditionally the highest priority.

Scheduling: Null Process

What does the CPU do when it has nothing to do (i.e., all the ready queue(s) are empty)?

Scheduling: Null Process

What does the CPU do when it has nothing to do (i.e., all the ready queue(s) are empty)? NULL...

Scheduling: Null Process

What does the CPU do when it has nothing to do (i.e., all the ready queue(s) are empty)? NULL...

- ▶ CPU has to do something, unfortunately it cannot do nothing
 - ▶ Simple solution: NULL process
 - ▶ The NULL process should have the lowest priority possible (What will happen if you give it the highest?)
 - ▶ Example NULL process:

```
void null_process() {  
    while (1) {  
        release_processor();  
    }  
}
```

Scheduling: `release_processor()`

1. Set `current_process` to state ready
2. `rpq_enqueue(current_process)` put `current_process` in ready queues
3. `process_switch` invokes scheduler and context-switches to the new process

Initialization

When the OS starts up, what initialization needs to happen?

- ▶ Initialize all hardware, including:
 - ▶ Serial port(s) and timer(s)
 - ▶ Memory mapping (memory allocation for memory blocks and stacks. . .)
 - ▶ Interrupts (hardware and software: vector table & traps)
- ▶ Create all kernel data structures
 - ▶ PCBs (status=ready), queues. . .
 - ▶ Place PCBs into respective queues

Initialization: Initialization Table

How does RTX know which processes to create?

Initialization: Initialization Table

How does RTX know which processes to create? You'll have a pre-defined initialization table.

Initialization: Initialization Table

How does RTX know which processes to create? You'll have a pre-defined initialization table.

- ▶ The table is an array of records
- ▶ Each record contains specifications of its process and any additional data structures required (unless you want to allocate those dynamically).

Example record contents:

Process ID
Priority
Initial SP
Initial PC

Questions

Do you have any?